

PROJECT DESCRIPTION

The project is an essential part of this class. It will allow you to demonstrate your Machine Learning (ML) skills and create something that you are proud of. It can also be a valuable addition to your projects portfolio that you can demonstrate to prospective employers.

Project Requirements

For the project, you have to perform a thorough processing and analysis of a dataset using ML techniques. Some of the requirements of the project are:

- The datasets should be chosen from a standard repository, such as Kaggle competitions, KDD cup competitions or choose another dataset having sufficient size and complexity. If you are not sure, please consult the instructor or the TA.
- You should apply multiple techniques and algorithms to the same dataset, and also compare their performance. In the end, you should identify which is your strongest technique and use that as your competition entry.
- You should use "strong" or "powerful" learners. Examples could be:
 - Deep Learning techniques
 - Ensemble Learning techniques, for example boosting or random forests
 - SVM with non-linear kernels
 - Recent ML libraries such as
 - Spark MLlib: <http://spark.apache.org/docs/latest/mllib-guide.html>
 - Flink: <https://flink.apache.org/news/2015/06/24/announcing-apache-flink-0.9.0-release.html>
 - Storm: <http://storm.apache.org/>
 - GO language: <http://www.datasciencecentral.com/profiles/blogs/machine-learning-libraries-in-go-language-3>
- Your results should be strong enough in terms of accuracy and other evaluation metrics e.g. ROC curve, area under ROC curve, and this will be one of the criteria for grades. Note that just using accuracy as the evaluation criteria is not sufficient.
- You should create a well formatted project **report** that should cover the following sections:
 - Introduction and problem description,
 - Related work
 - Dataset description (including features, attributes, etc)
 - Pre-processing techniques

- Your proposed solution, and methods [This section should have enough details – both theoretical, and practical]
- Experimental results and analysis [Details are expected]
- Conclusion
- Contribution of team members
- References

An excellent example of what to include in such a report can be found here:
<http://www.cs.utexas.edu/~mooney/cs391L/paper-template.html>

Some examples of excellent reports can be found at: (Note: You cannot choose these project topics)

<http://cs229.stanford.edu/projects2015.html>

<http://cs229.stanford.edu/projects2014.html>

<http://cs229.stanford.edu/projects2013.html>

All contents of your report must be original. You cannot copy sentences, paragraphs, figures, or anything else from outside sources. As a graduate student, you are expected to work with maturity and diligence.

Again, your report will be checked for plagiarism. Any violation will carry strong penalties, including reporting the incident to university authorities.

- Team size requirements: Project can be done in teams of 1 to 4 students. More than 4 students cannot be in a team under any circumstances. You can only form team within the same class and section. You are not allowed to work or collaborate with students from other sections of this class.
- Project selections should be approved by the instructor. It is permissible for more than one team to work on the same topic, if it's approved by the instructor.
- After selecting your project, please be sure to fill out your details here:
<https://goo.gl/forms/TFJE6JlBeXq3d91x2>
- The final project report is due at midnight Friday Nov 30. Project demos and presentations will be required in front of the TA during the last week of the semester. These are strict deadlines.

Project Ideas

Below are some of the project ideas. You can choose any one of them. Note that for the data science competitions, you have multiple options. You are free to choose any active competition, but you will have to follow the requirements completely. You are expected to meet as many requirements of the competition as possible.

Note: Before starting to work on your project, make sure you have the instructor's approval which will be granted on the Google spreadsheet.

1. Participate in the Yelp dataset challenge and submit a good entry:

http://www.yelp.com/dataset_challenge

2. Take part in an **active** Kaggle competition that involves significant amount of Machine Learning technologies

<https://www.kaggle.com/competitions>

3. Take part in KDD cup challenge

<http://www.kdd.org/kdd-cup>

You can take part in any previous year's cup.

4. Take part in an **active** Driven Data competition.

<https://www.drivendata.org/>

5. Machine learning based analysis of stock market investing techniques

Ideas:

- Simulation of systematic trading techniques, such as backtesting
https://en.wikipedia.org/wiki/Technical_analysis#Systematic_trading
- Simulation and analysis of backtesting using R packages such as backtest, PerformanceAnalytics, quantmod, etc

6. Take part in a competition from KDnuggets

<https://www.kdnuggets.com/competitions/>

7. Take part in a competition from Innocentive

<https://www.innocentive.com/ar/challenge/browse>

8. Take part in a competition from Tunedit
<http://tunedit.org/>

9. Take part in a TopCode challenge:
<https://www.topcoder.com/community/data-science/>

10. You are also free to propose your own topic, which could involve using Deep Learning techniques, such as Convolution or Recurrent Neural Networks, etc.

Deliverables and Deadlines

Deadline	Project Phase	Deliverable
Friday Oct 26 Midnight	Project Selection Team Formation	Submit your details on Google Forms https://goo.gl/forms/TFJE6JlBeXq3d91x2 Please check for instructor's comments and approval at: https://goo.gl/CZpmb9
Friday Nov 9 Midnight	Project Status Report	Submit a report containing following on eLearning: <ul style="list-style-type: none">• Dataset details, such as number of features, instances, data distribution• Techniques you plan to use• Experimental methodology (how you plan to pre-process, create training, validation, and test datasets, and other such details)• Coding language / technique to be used• Preliminary Results (if available)
Friday Nov 30 Midnight	Final Report	Submit final documents on eLearning: <ul style="list-style-type: none">• Detailed Final Project Report• Code• README file indicating how to run your code ** Your report and code will be checked for plagiarism **